

# Bit Vector Decision Procedures

*A Basis for Reasoning about  
Hardware & Software*

Randal E. Bryant  
Carnegie Mellon University

<http://www.cs.cmu.edu/~bryant>

# Collaborators

## Sanjit Seshia, Bryan Brady

- UC Berkeley

## Daniel Kroening

- ETH Zurich

## Joel Ouaknine

- Oxford University

## Ofer Strichman

- Technion

## Randy Bryant

- Carnegie Mellon

3 continents  
4 countries  
5 institutions  
6 authors

# Motivating Example #1

```
int abs(int x) {  
    int mask = x>>31;  
    return (x ^ mask) + ~mask + 1;  
}
```

```
int test_abs(int x) {  
    return (x < 0) ? -x : x;  
}
```

***Do these functions produce identical results?***

## Strategy

- Represent and reason about bit-level program behavior
- Specific to machine word size, integer representations, and operations

# Motivating Example #2

```
void fun() {  
    char fmt[16];  
    fgets(fmt, 16, stdin);  
    fmt[15] = '\0';  
    printf(fmt);  
}
```

*Is there an input string that causes value 234 to be written to address  $a_4a_3a_2a_1$ ?*

## Answer

- Yes: " $a_1a_2a_3a_4\%230g\%n$ "
  - Depends on details of compilation
- But no exploit for buffer size less than 8
- [Ganapathy, Seshia, Jha, Reps, Bryant, ICSE '05]

# Motivating Example #3

```
bit[W] popSpec(bit[W] x)
{
    int cnt = 0;
    for (int i=0; i<W; i++) {
        if (x[i]) cnt++;
    }
    return cnt;
}
```

```
bit[W] popSketch(bit[W] x)
{
    loop (??) {
        x = (x&??) + ((x>>??)&??);
    }
    return x;
}
```

*Is there a way to expand the program sketch to make it match the spec?*

## Answer

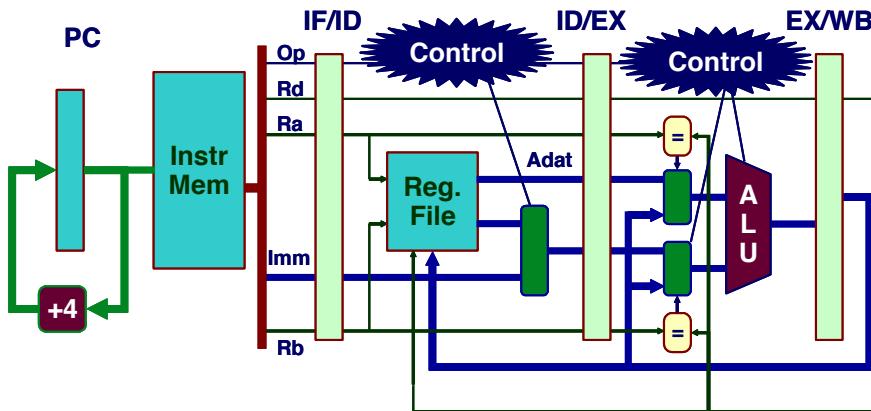
- W=16:

```
x = (x&0x5555) + ((x>>1)&0x5555);
x = (x&0x3333) + ((x>>2)&0x3333);
x = (x&0x0077) + ((x>>8)&0x0077);
x = (x&0x000f) + ((x>>4)&0x000f);
```

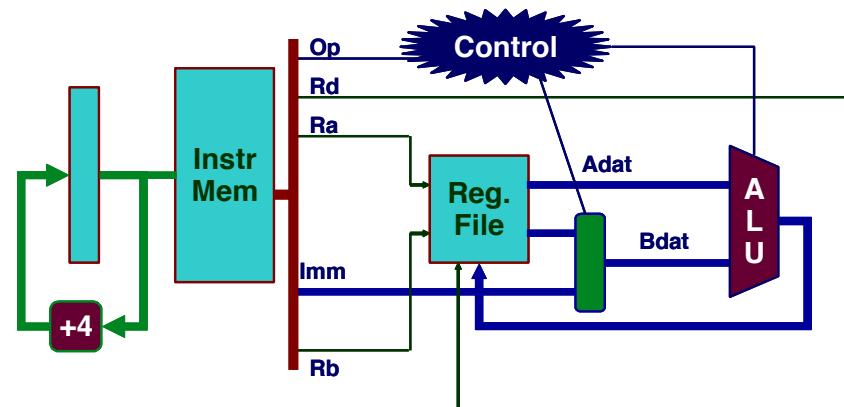
- [Solar-Lezama, et al., ASPLOS '06]

# Motivating Example #4

Pipelined Microprocessor



Sequential Reference Model



*Is pipelined microprocessor identical to sequential reference model?*

## Strategy

- Automatically generate abstraction function from pipeline to program state [Burch & Dill, CAV '94]
- Represent machine instructions, data, and state as bit vectors
  - Compatible with hardware description language representation

# Task

## Bit Vector Formulas

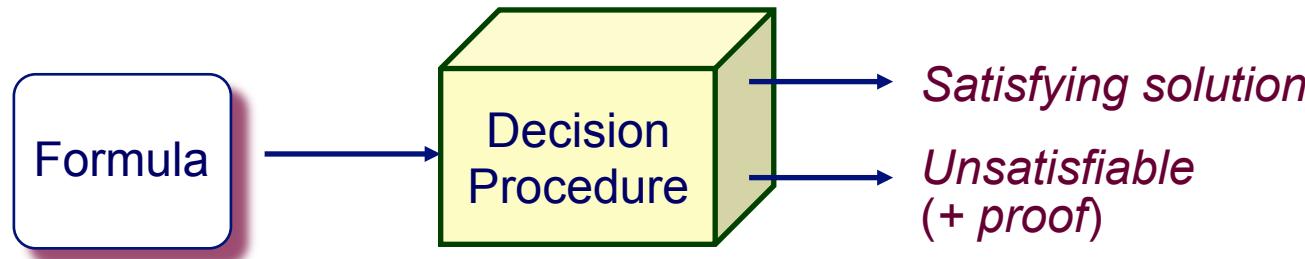
- Fixed width data words
- Arithmetic operations
  - E.g., add/subtract/multiply/divide & comparisons
  - Two's complement, unsigned, ...
- Bit-wise logical operations
  - E.g., and/or/xor, shift/extract and equality
- Boolean connectives

## Reason About Hardware & Software at Bit Level

- Formal verification
- Security analysis
- Test & program generation
  - What function arguments will cause program to take specified branch?

# Decision Procedures

- Core technology for formal reasoning



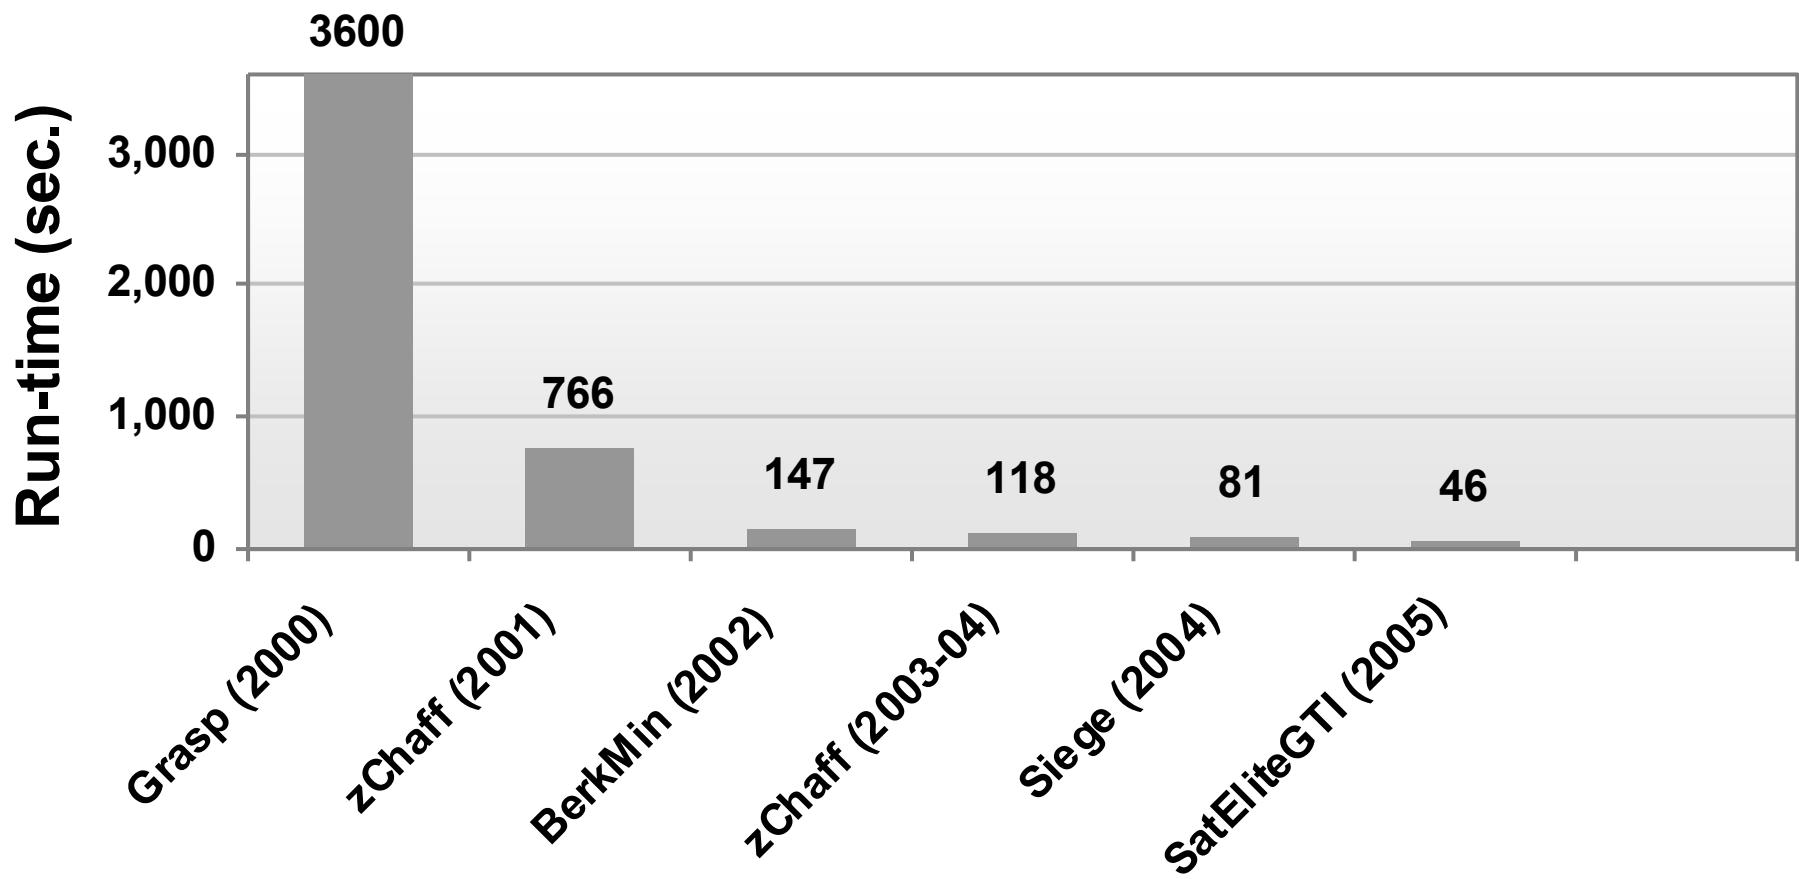
## Boolean SAT

- Pure Boolean formula

## SAT Modulo Theories (SMT)

- Support additional logic fragments
- Example theories
  - Linear arithmetic over reals or integers
  - Functions with equality
  - *Bit vectors*
  - Combinations of theories

# Recent Progress in SAT Solving



# BV Decision Procedures: Some History

## B.C. (Before Chaff)

- String operations (concatenate, field extraction)
- Linear arithmetic with bounds checking
- Modular arithmetic

## SAT-Based “Bit Blasting”

- Generate Boolean circuit based on bit-level behavior of operations
- Convert to Conjunctive Normal Form (CNF) and check with best available SAT checker
- Handles arbitrary operations
- Effective in many applications
  - CBMC [Clarke, Kroening, Lerda, TACAS '04]
  - Microsoft Cogent + SLAM [Cook, Kroening, Sharygina, CAV '05]
  - CVC-Lite [Dill, Barrett, Ganesh], Yices [deMoura, et al], STP

# Research Challenge

*Is there a better way than bit blasting?*

## Requirements

- Provide same functionality as with bit blasting
- Find abstractions based on word-level structure
- Improve on performance of bit blasting

## A New Approach

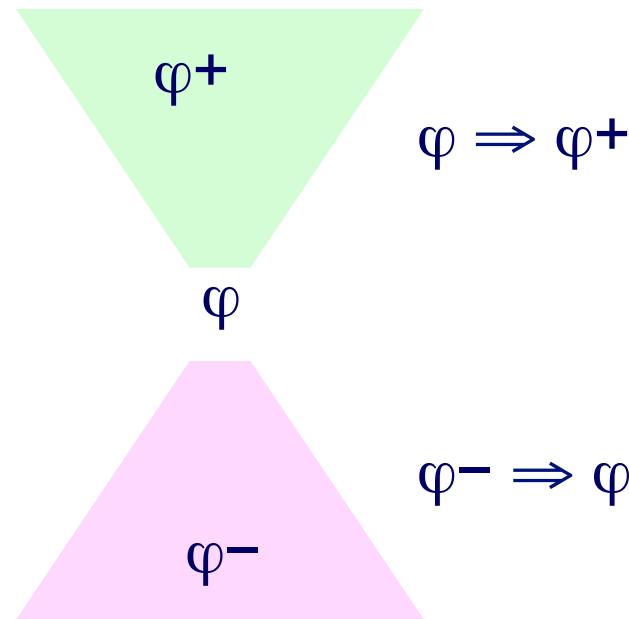
- [Bryant, Kroening, Ouaknine, Seshia, Stichman, Brady, TACAS '07]
- Use bit blasting as core technique
- Apply to simplified versions of formula
- Successive approximations until solve or show unsatisfiable

# Approximating Formula

Overapproximation

Original Formula

Underapproximation



**More solutions:**  
If unsatisfiable,  
then so is  $\varphi$

**Fewer solutions:**  
Satisfying solution  
also satisfies  $\varphi$

## Example Approximation Techniques

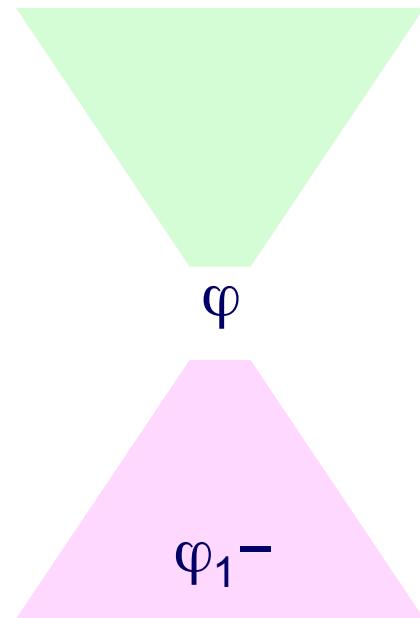
### ■ Underapproximating

- Restrict word-level variables to smaller ranges of values

### ■ Overapproximating

- Replace subformula with Boolean variable

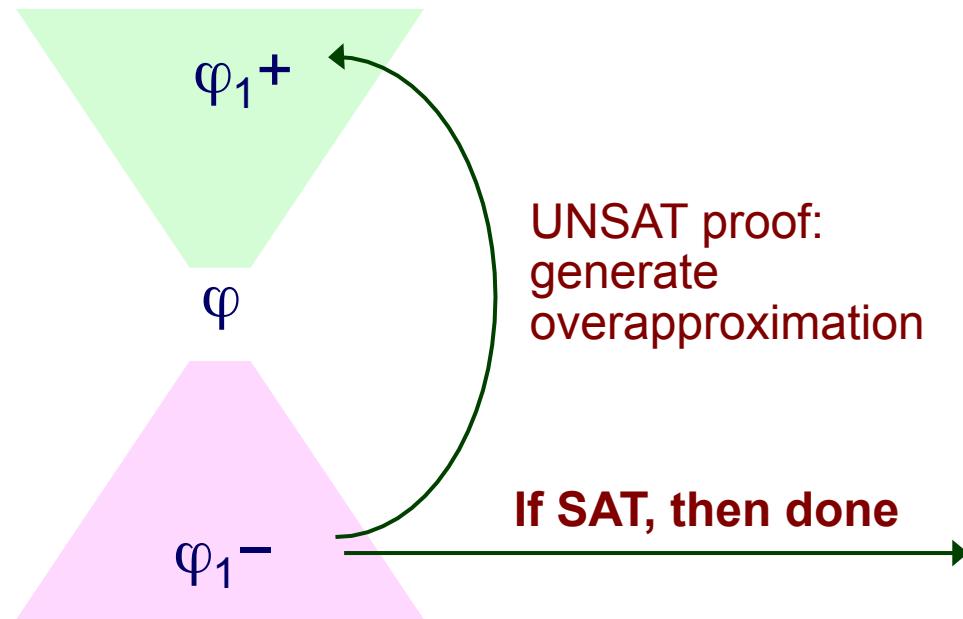
# Starting Iterations



## Initial Underapproximation

- (Greatly) restrict ranges of word-level variables
- Intuition: Satisfiable formula often has small-domain solution

# First Half of Iteration

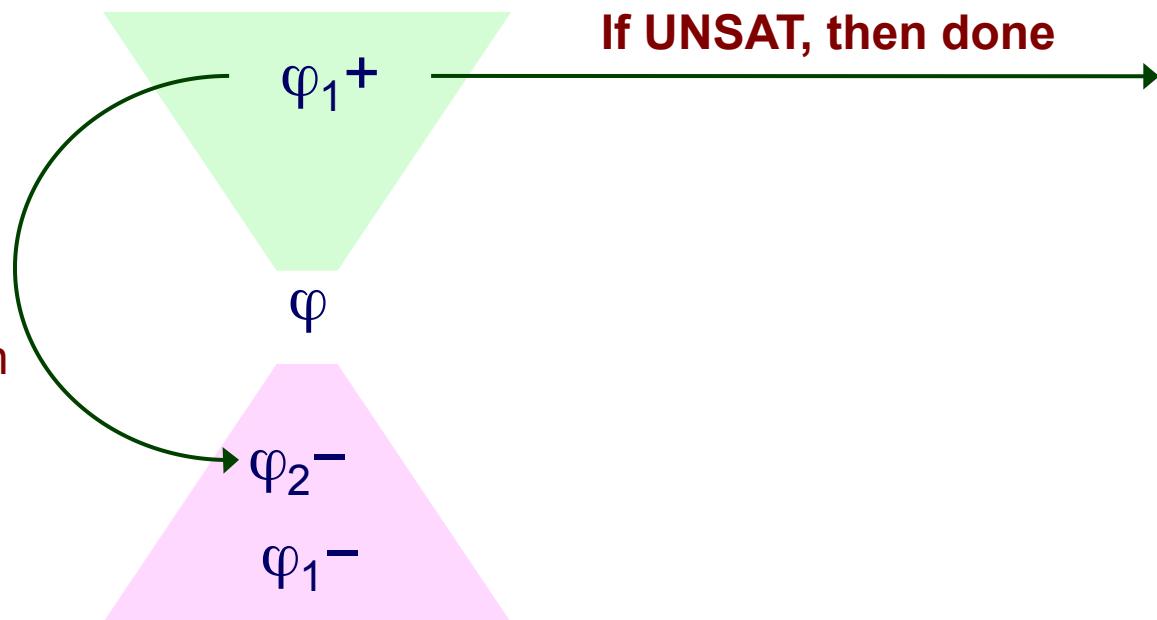


## SAT Result for $\varphi_1^-$

- **Satisfiable**
  - Then have found solution for  $\varphi$
- **Unsatisfiable**
  - Use UNSAT proof to generate overapproximation  $\varphi_1^+$
  - (Described later)

# Second Half of Iteration

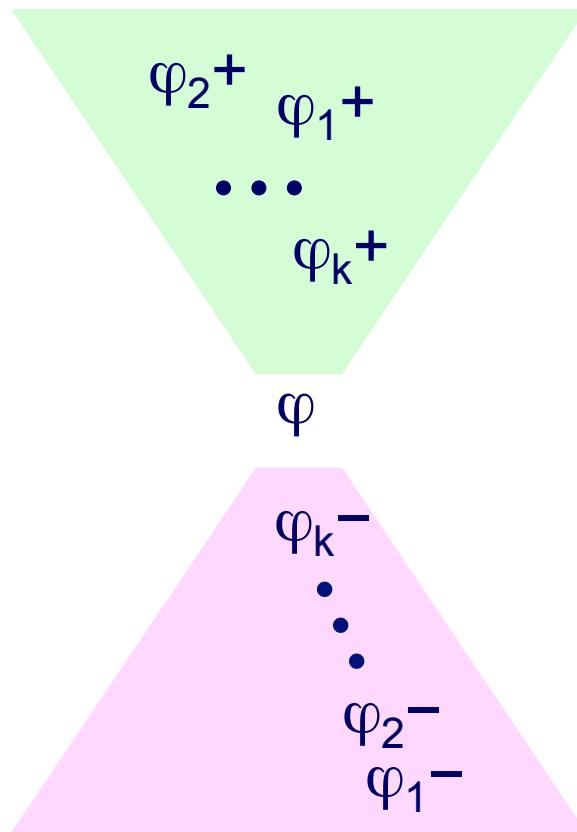
SAT:  
Use solution to generate  
refined underapproximation



## SAT Result for $\varphi_1^+$

- **Unsatisfiable**
  - Then have shown  $\varphi$  unsatisfiable
- **Satisfiable**
  - Solution indicates variable ranges that must be expanded
  - Generate refined underapproximation

# Iterative Behavior



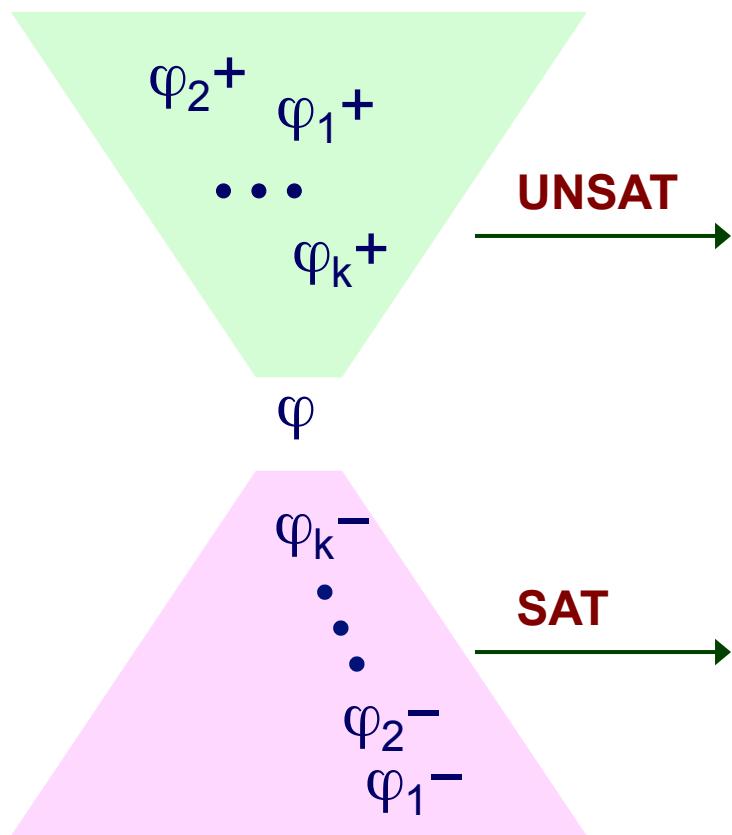
## Underapproximations

- Successively more precise abstractions of  $\varphi$
- Allow wider variable ranges

## Overapproximations

- No predictable relation
- UNSAT proof not unique

# Overall Effect



## Soundness

- Only terminate with solution on underapproximation
- Only terminate as UNSAT on overapproximation

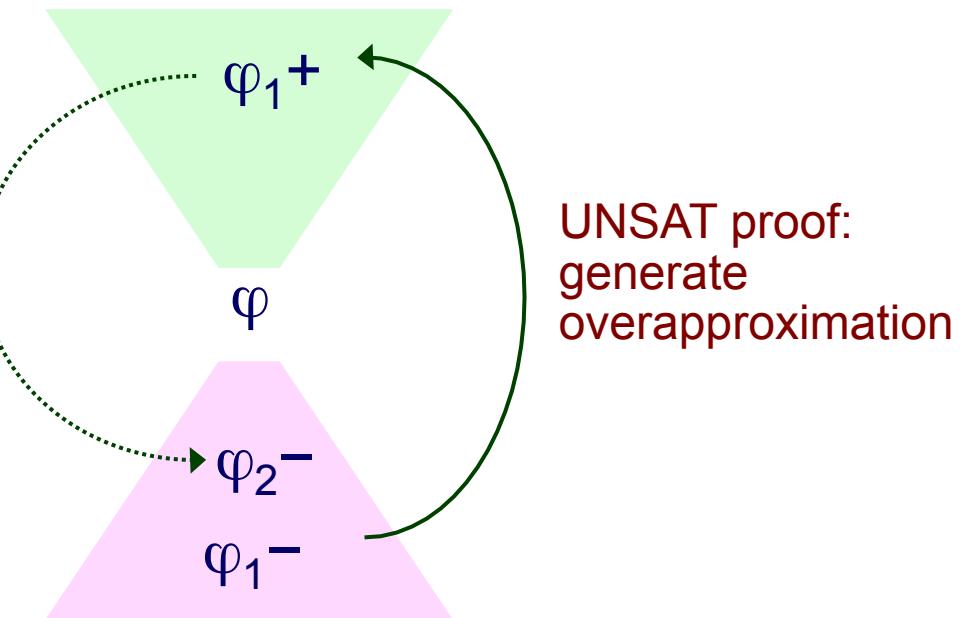
## Completeness

- Successive underapproximations approach  $\varphi$
- Finite variable ranges guarantee termination
  - In worst case, get  $\varphi_k^- = \varphi$

# Generating Overapproximation

## Given

- Underapproximation  $\varphi_1^-$
- Bit-blasted translation of  $\varphi_1^-$  into Boolean formula
- Proof that Boolean formula unsatisfiable

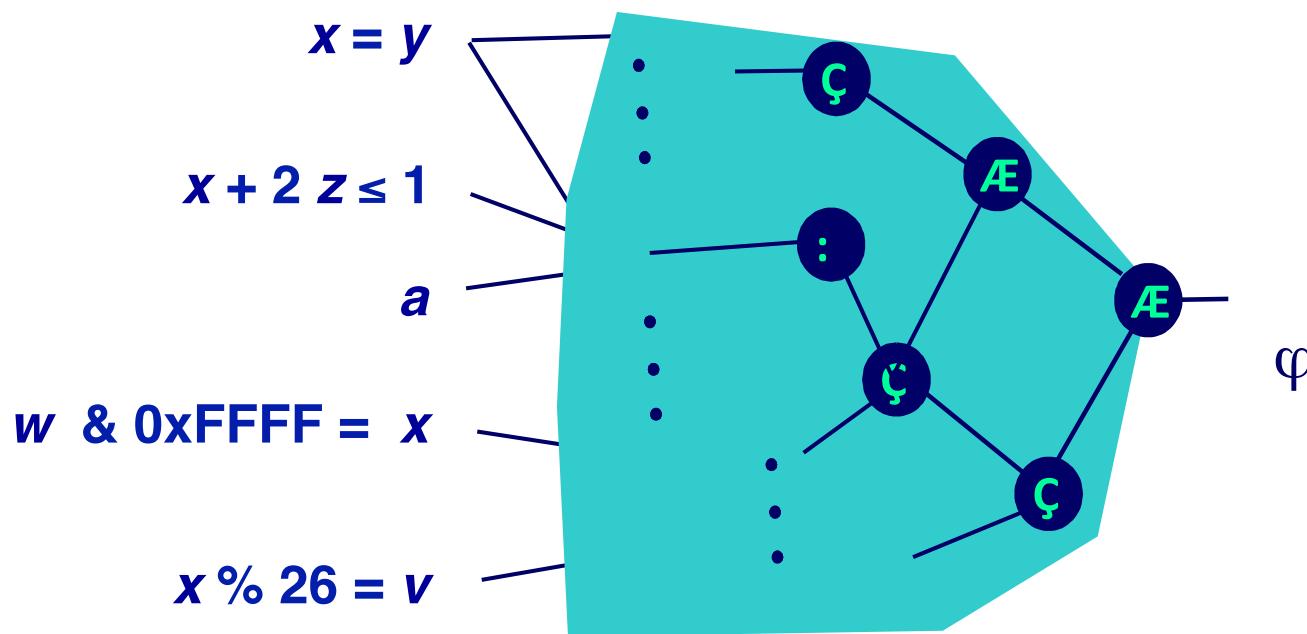


## Generate

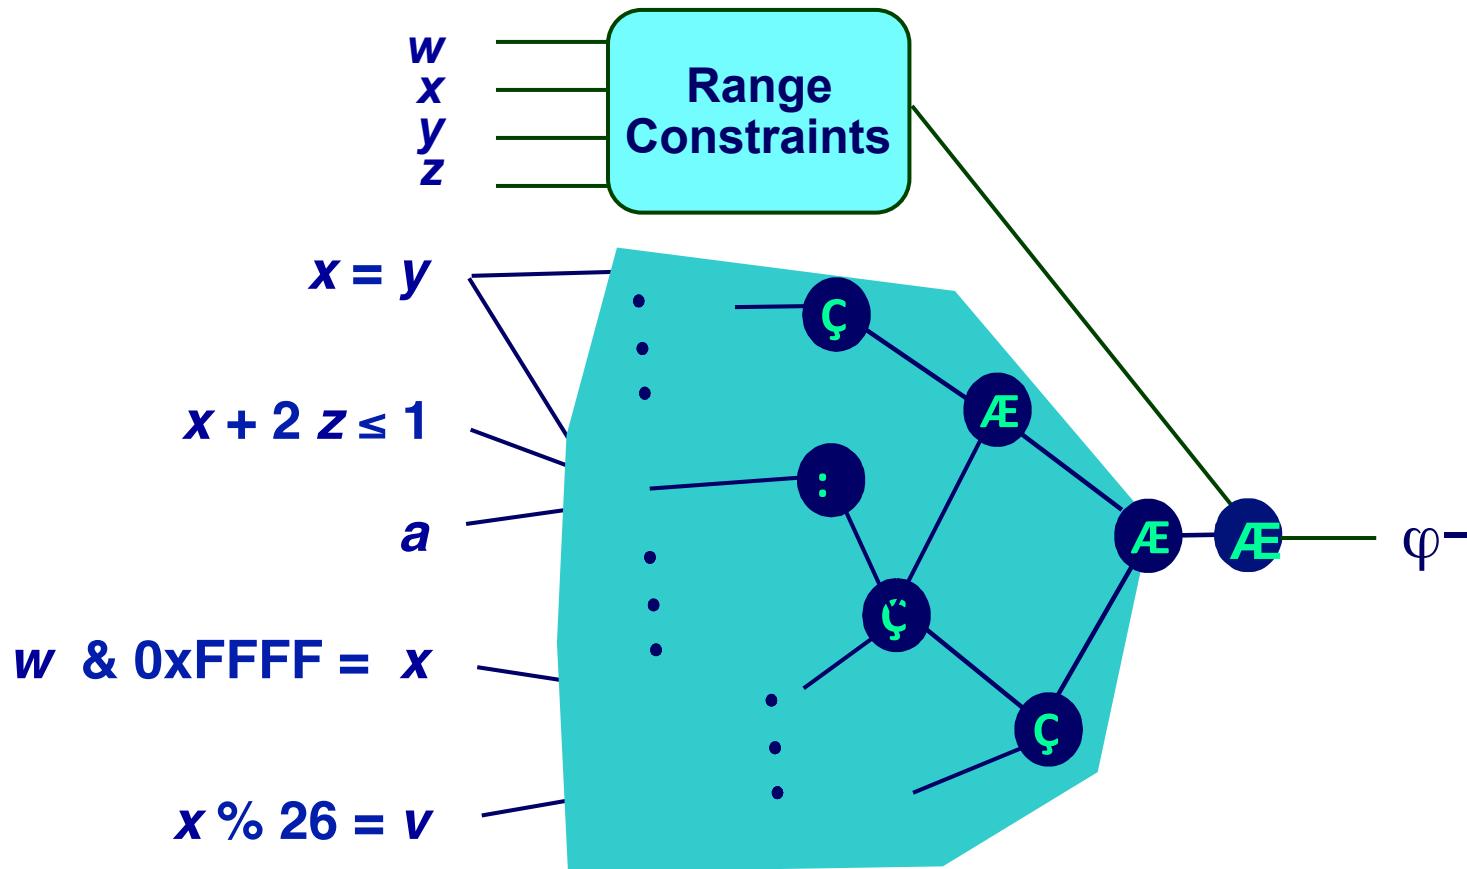
- Overapproximation  $\varphi_1^+$
- If  $\varphi_1^+$  satisfiable, must lead to refined underapproximation
  - Generate  $\varphi_2^-$  such that  $\varphi_1^- \Rightarrow \varphi_2^- \Rightarrow \varphi$

# Bit-Vector Formula Structure

- DAG representation to allow shared subformulas



# Structure of Underapproximation



- **Linear complexity translation to CNF**
  - Each word-level variable encoded as set of Boolean variables
  - Additional Boolean variables represent subformula values

# Encoding Range Constraints

## Explicit

- View as additional predicates in formula



## Implicit

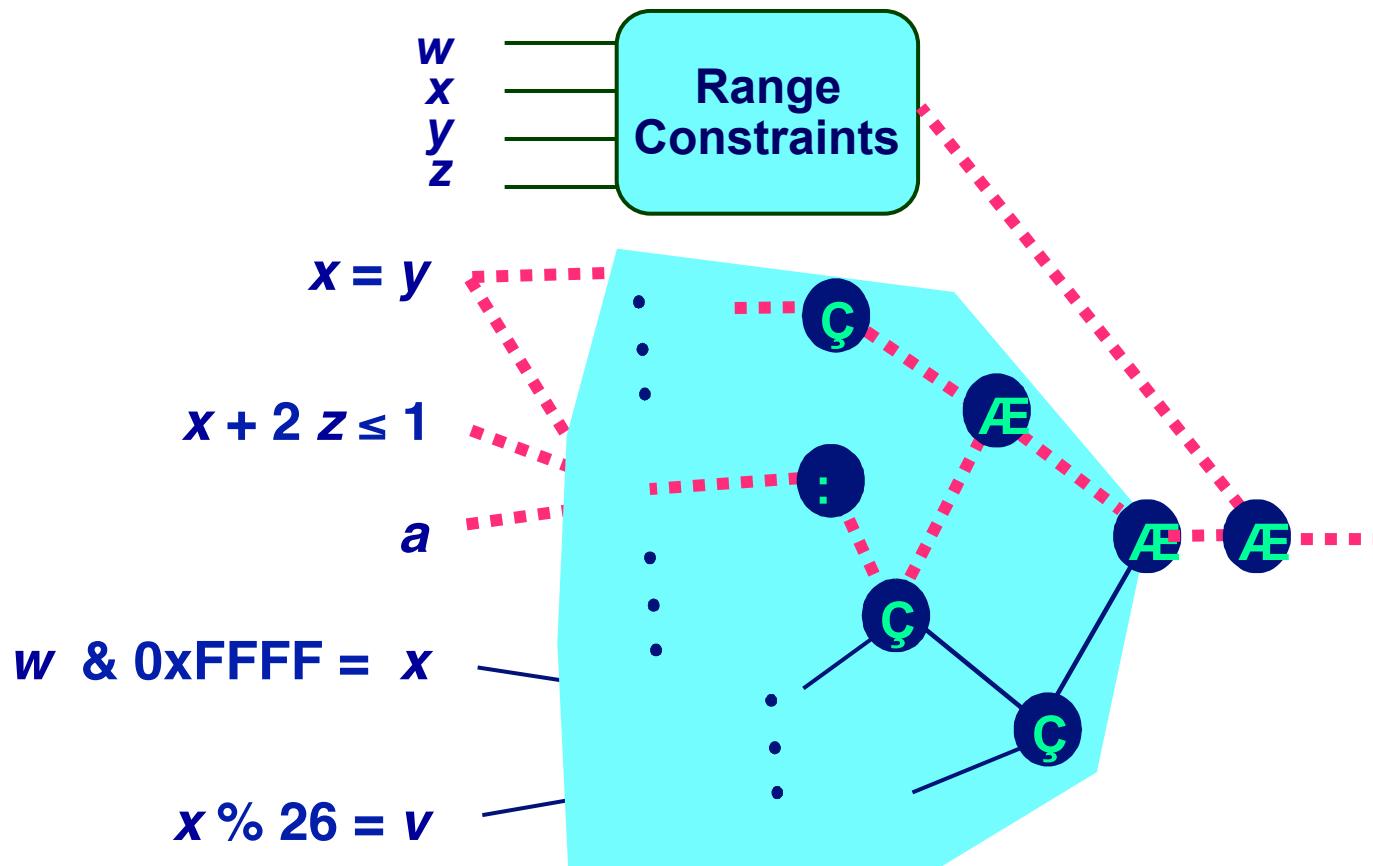
- Reduce number of variables in encoding

Constraint	Encoding
$0 \leq w < 8$	$0\ 0\ 0 \cdots 0\ w_2 w_1 w_0$
$-4 \leq x < 4$	$x_s x_s x_s \cdots x_s x_s x_1 x_0$

- Yields smaller SAT encodings

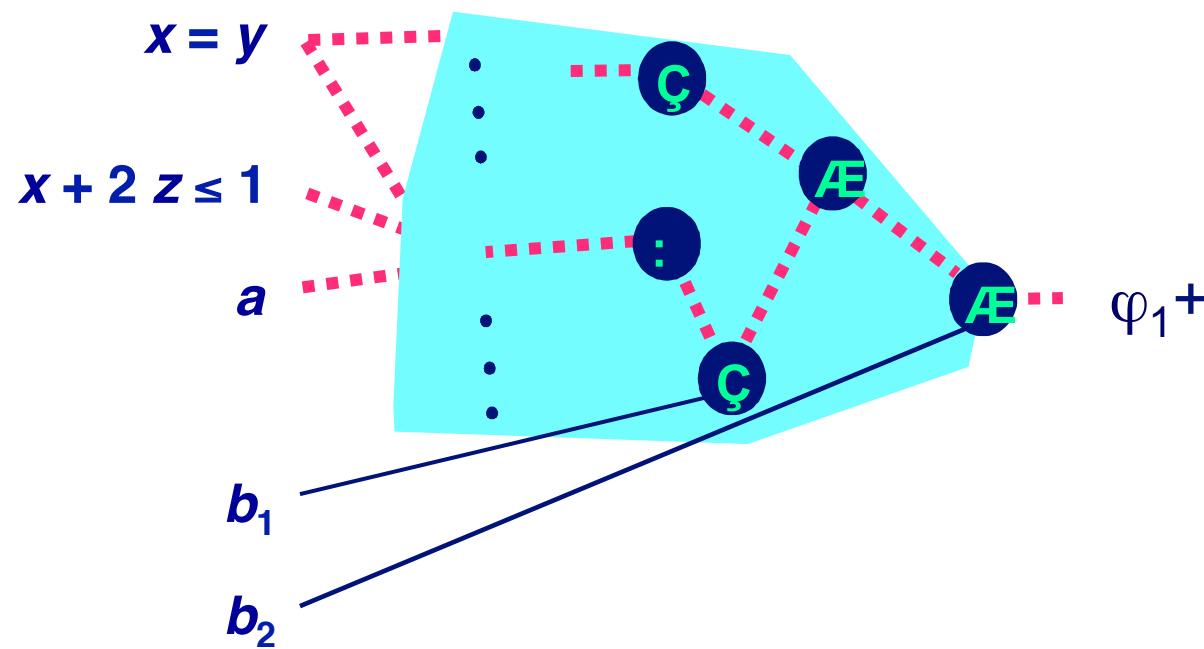
# UNSAT Proof

- Subset of clauses that is unsatisfiable
- Clause variables define portion of DAG
- Subgraph that cannot be satisfied with given range constraints



# Generated Overapproximation

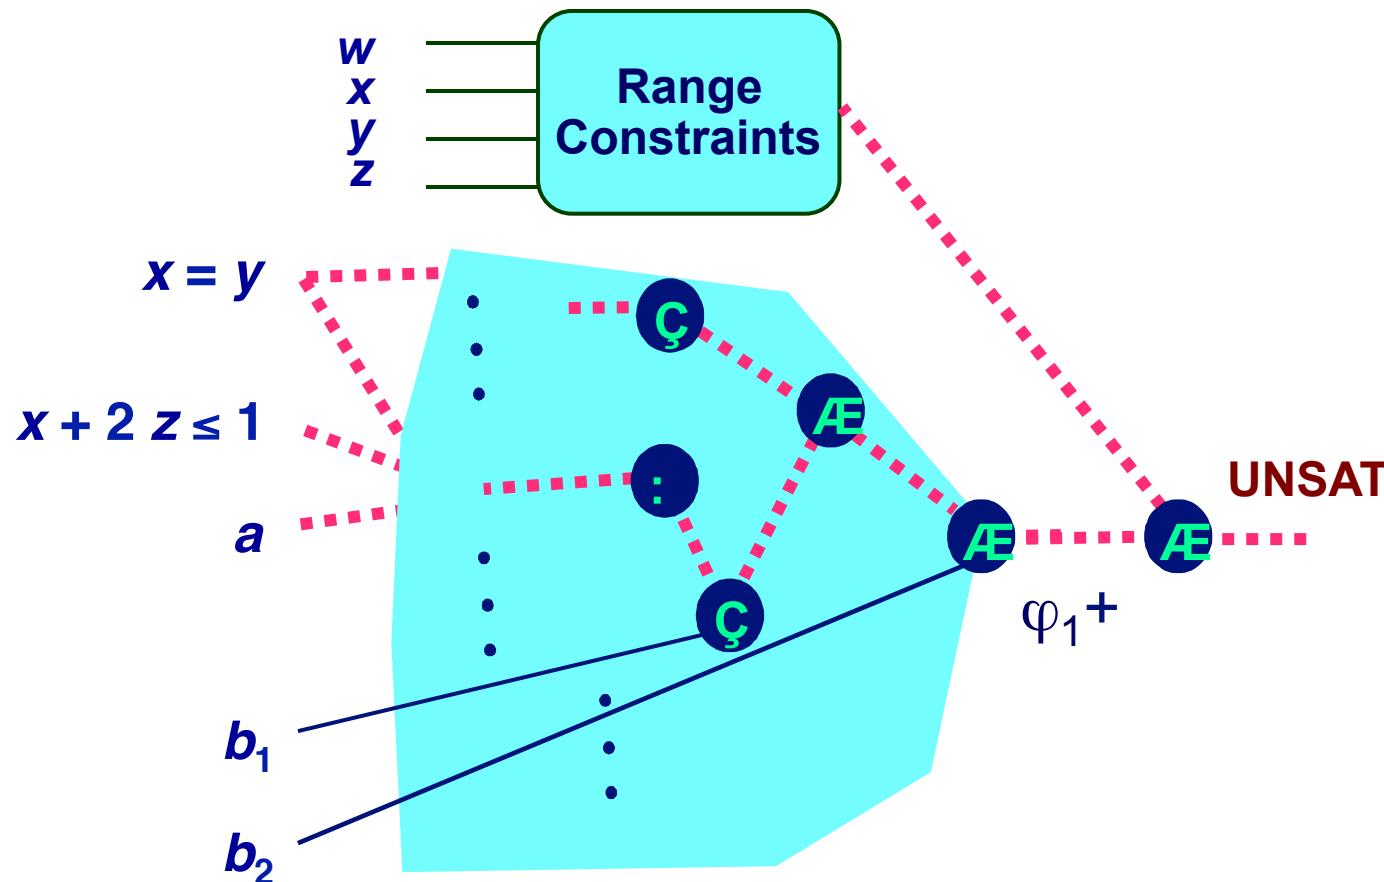
- Identify subformulas containing no variables from UNSAT proof
- Replace by fresh Boolean variables
- Remove range constraints on word-level variables
- Creates overapproximation
  - Ignores correlations between values of subformulas



# Refinement Property

## Claim

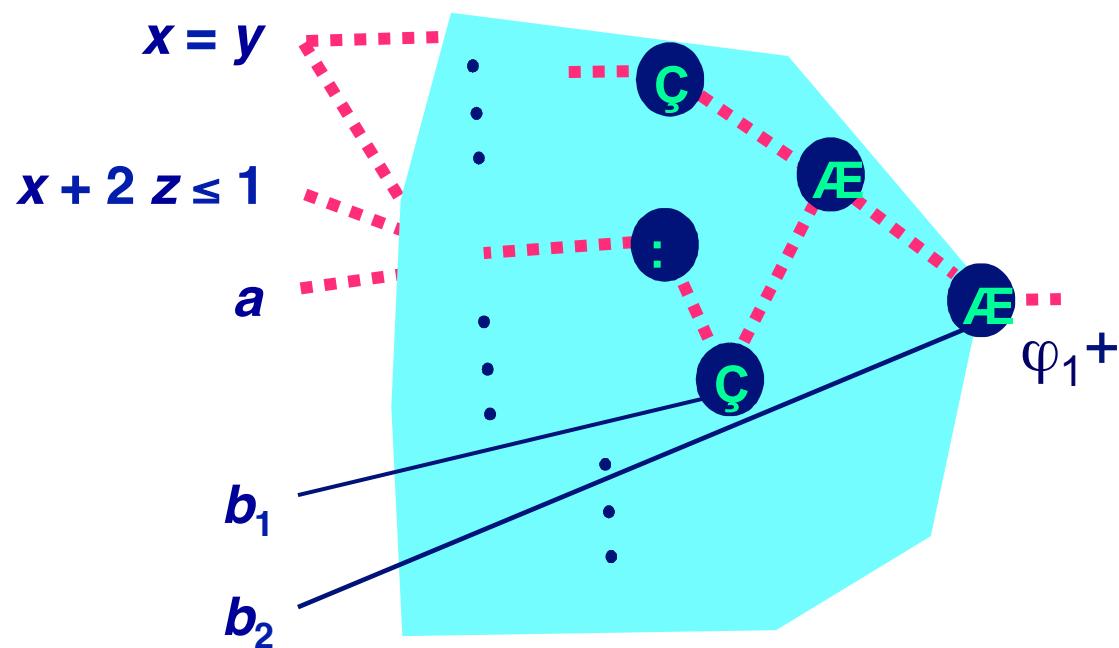
- $\varphi_1^+$  has no solutions that satisfy  $\varphi_1^-$ 's range constraints
  - Because  $\varphi_1^+$  contains portion of  $\varphi_1^-$  that was shown to be unsatisfiable under range constraints



# Refinement Property (Cont.)

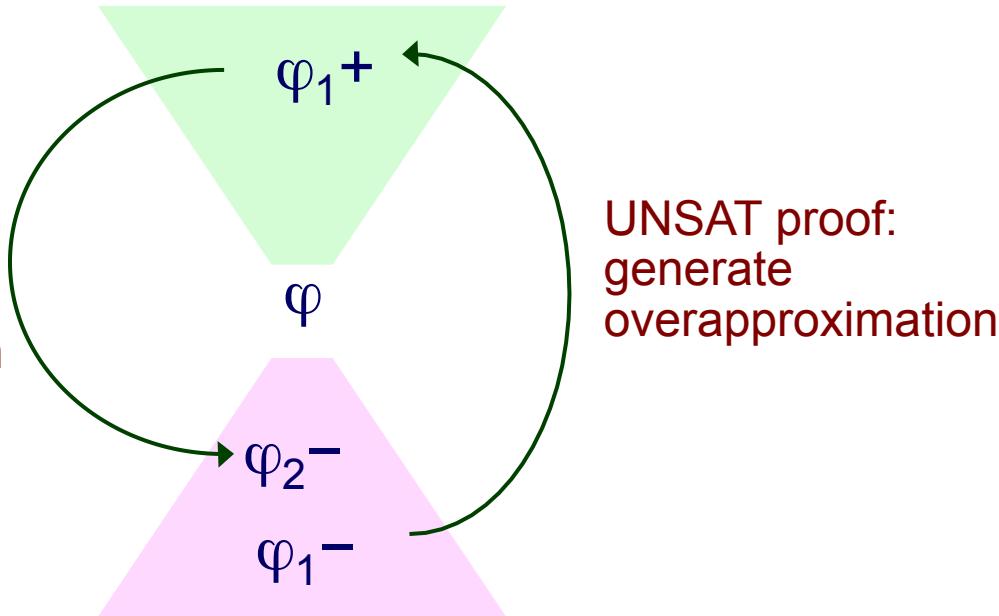
## Consequence

- Solving  $\varphi_1+$  will expand range of some variables
- Leading to more exact underapproximation  $\varphi_2-$



# Effect of Iteration

SAT:  
Use solution to generate  
refined underapproximation



UNSAT proof:  
generate  
overapproximation

## Each Complete Iteration

- Expands ranges of some word-level variables
- Creates refined underapproximation

# Approximation Methods

## So Far

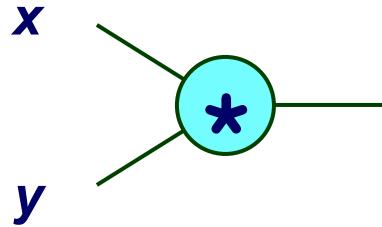
- **Range constraints**
  - Underapproximate by constraining values of word-level variables
- **Subformula elimination**
  - Overapproximate by assuming subformula value arbitrary

## General Requirements

- **Systematic under- and over-approximations**
- **Way to connect from one to another**

## Goal: Devise Additional Approximation Strategies

# Function Approximation Example



		x		
		0	1	else
y	0	0	0	0
	1	0	1	x
	else	0	y	§

## Motivation

- Multiplication (and division) are difficult cases for SAT

## §: Prohibited

- Gives underapproximation
- Restricts values of (possibly intermediate) terms

## §: $f(x,y)$

- Overapproximate as uninterpreted function  $f$
- Value constrained only by functional consistency

# Results: UCLID BV vs. Bit-blasting

Formula	Ans.	Bit-Blasting			UCLID			STP (sec.)	Yices (sec.)		
		Run-time (sec.)			Run-time (sec.)						
		Enc.	SAT	Total	Enc.	SAT	Total				
Y86-std	UNSAT	17.91	TO	TO	23.51	987.91	<b>1011.42</b>	2083.73	TO		
Y86-btnft	UNSAT	17.79	TO	TO	26.15	1164.07	<b>1190.22</b>	err	TO		
s-40-50	SAT	6.00	33.46	39.46	106.32	10.45	116.77	<b>12.96</b>	65.51		
BBB-32	SAT	37.09	29.98	67.07	19.91	1.74	<b>21.65</b>	38.45	183.30		
runit_flat-64	SAT	121.99	32.16	154.15	19.52	1.68	<b>21.20</b>	873.67	1312.00		
C1-P1	SAT	2.68	45.19	47.87	2.61	0.58	<b>3.19</b>	err	err		
C1-P2	UNSAT	0.44	TO	TO	2.24	2.12	<b>4.36</b>	TO	TO		
C3-OP80	SAT	14.96	TO	TO	14.54	349.41	<b>363.95</b>	TO	3242.43		
egt-5212	UNSAT	0.064	0.003	0.067	0.163	0.001	0.164	0.018	<b>0.009</b>		

[results on 2.8 GHz Xeon, 2 GB RAM]

- **UCLID always better than bit blasting**
- **Generally better than other available procedures**
- **SAT time is the dominating factor**

# UCLID BV run-time analysis

Formula	Ans.	$\max_i s_i$	$\max_i w_i$	Num. Iter	$\max \frac{ \bar{\phi} }{ \phi }$	Speedup
Y86-std	UNSAT	4	32	1	0.18	2.06
Y86-btnft	UNSAT	4	32	1	0.20	> 3.01
s-40-50	SAT	32	32	8	0.12	0.11
BBB-32	SAT	4	32	1	—	1.78
rfunit_flat-64	SAT	4	64	1	—	7.27
C1-P1	SAT	2	65	1	—	15.00
C1-P2	UNSAT	2	14	1	1.00	> 825.69
C3-OP80	SAT	2	9	1	—	8.91
egt-5212	UNSAT	8	8	1	0.13	0.06

- **$w_i$ : Maximum word-level variable size**
- **$s_i$ : Maximum word-level variable instantiation**
- **Generated abstractions are small**
- **Few iterations of refinement loop needed**

# Why This Work is Worthwhile

## Realistic Semantic Model for Hardware and Software

- Captures all details of actual operation
  - Detects errors related to overflow and other artifacts of finite representation
- Allows mixing of integer and bit-level operations
- Can capture many abstractions that are currently applied manually

## SAT-Based Methods Are Only Logical Choice

- Bit blasting is only way to capture full set of operations
- SAT solvers are good & getting better

## Abstraction / Refinement Allows Better Scaling

- Take advantage of cases where formula easily satisfied or disproven

# Future Work

## Lots of Refinement & Tuning

- Selecting under- and over-approximations
- Iterating within under- or over-approximation
  - E.g., attempt to control variable ranges when overapproximating
- Reusing portions of bit-blasted formulas
  - Take advantage of incremental SAT

## Additional Abstractions

- More general use of functional abstraction
  - Subsume use of uninterpreted functions in current verification methods